

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as py
import seaborn as sns
```

```
In [2]: d=pd.read_csv(r"C:\Users\user\Downloads\fiat500_VehicleSelection_Dataset (2).csv")
d
```

```
Out[2]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700
...	...	...	...	...	...	...	...	...	...
1533	1534	sport	51	3712	115280	1	45.069679	7.704920	5200
1534	1535	lounge	74	3835	112000	1	45.845692	8.666870	4600
1535	1536	pop	51	2223	60457	1	45.481541	9.413480	7500
1536	1537	lounge	51	2557	80750	1	45.000702	7.682270	5990
1537	1538	pop	51	1766	54276	1	40.323410	17.568270	7900

1538 rows × 9 columns

```
In [3]: d.head()
```

```
Out[3]:
```

	ID	model	engine_power	age_in_days	km	previous_owners	lat	lon	price
0	1	lounge	51	882	25000	1	44.907242	8.611560	8900
1	2	pop	51	1186	32500	1	45.666359	12.241890	8800
2	3	sport	74	4658	142228	1	45.503300	11.417840	4200
3	4	lounge	51	2739	160000	1	40.633171	17.634609	6000
4	5	pop	73	3074	106880	1	41.903221	12.495650	5700

```
In [4]: d.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 9 columns):
#   Column          Non-Null Count  Dtype
---  -
0   ID              1538 non-null   int64
1   model          1538 non-null   object
```

```

2  engine_power      1538 non-null  int64
3  age_in_days       1538 non-null  int64
4  km                1538 non-null  int64
5  previous_owners   1538 non-null  int64
6  lat               1538 non-null  float64
7  lon               1538 non-null  float64
8  price             1538 non-null  int64
dtypes: float64(2), int64(6), object(1)
memory usage: 108.3+ KB

```

```
In [5]: d.describe()
```

```
Out[5]:
```

	ID	engine_power	age_in_days	km	previous_owners	lat	lon
<b>count</b>	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000	1538.000000
<b>mean</b>	769.500000	51.904421	1650.980494	53396.011704	1.123537	43.541361	11.56341
<b>std</b>	444.126671	3.988023	1289.522278	40046.830723	0.416423	2.133518	2.32811
<b>min</b>	1.000000	51.000000	366.000000	1232.000000	1.000000	36.855839	7.24541
<b>25%</b>	385.250000	51.000000	670.000000	20006.250000	1.000000	41.802990	9.50501
<b>50%</b>	769.500000	51.000000	1035.000000	39031.000000	1.000000	44.394096	11.86921
<b>75%</b>	1153.750000	51.000000	2616.000000	79667.750000	1.000000	45.467960	12.76901
<b>max</b>	1538.000000	77.000000	4658.000000	235000.000000	4.000000	46.795612	18.36551

```
In [6]: d.columns
```

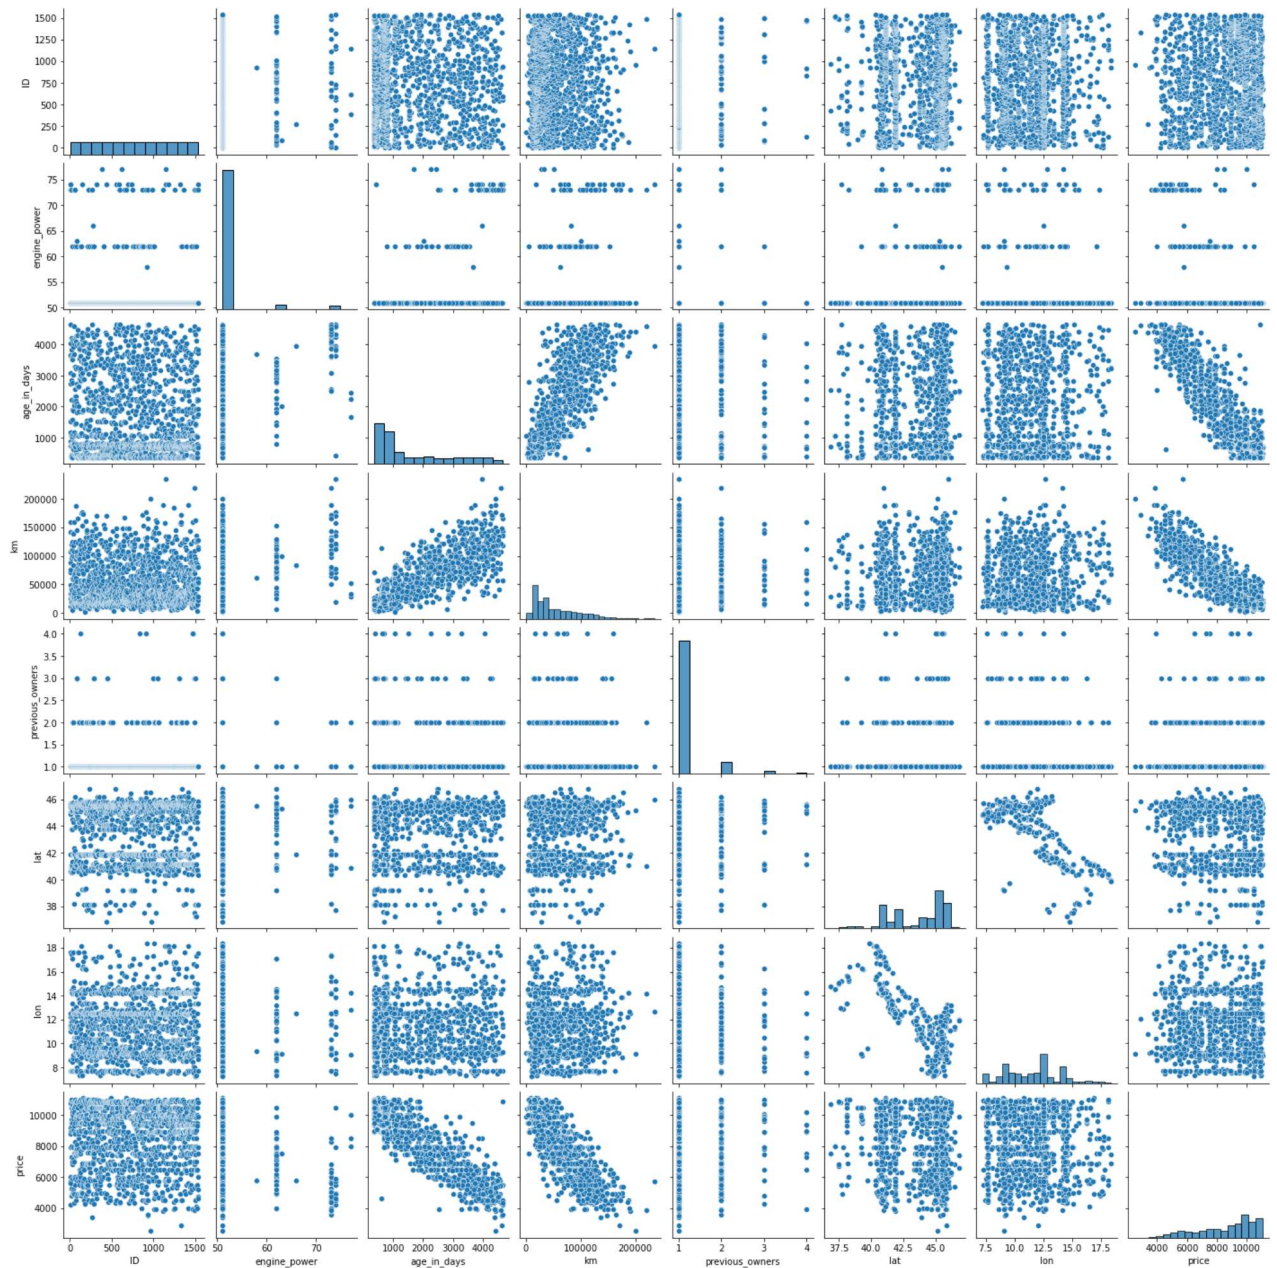
```
Out[6]: Index(['ID', 'model', 'engine_power', 'age_in_days', 'km', 'previous_owners',
              'lat', 'lon', 'price'],
              dtype='object')
```

```
In [7]: d.index
```

```
Out[7]: RangeIndex(start=0, stop=1538, step=1)
```

```
In [8]: sns.pairplot(d)
```

```
Out[8]: <seaborn.axisgrid.PairGrid at 0x1f3cbdaf4c0>
```

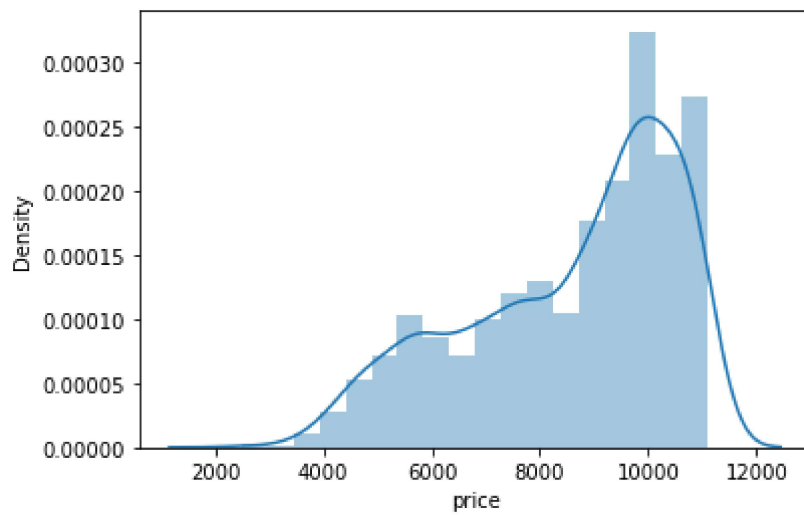


```
In [9]: sns.distplot(d['price'])
```

C:\ProgramData\Anaconda3\lib\site-packages\seaborn\distributions.py:2557: FutureWarning: `distplot` is a deprecated function and will be removed in a future version. Please adapt your code to use either `displot` (a figure-level function with similar flexibility) or `histplot` (an axes-level function for histograms).

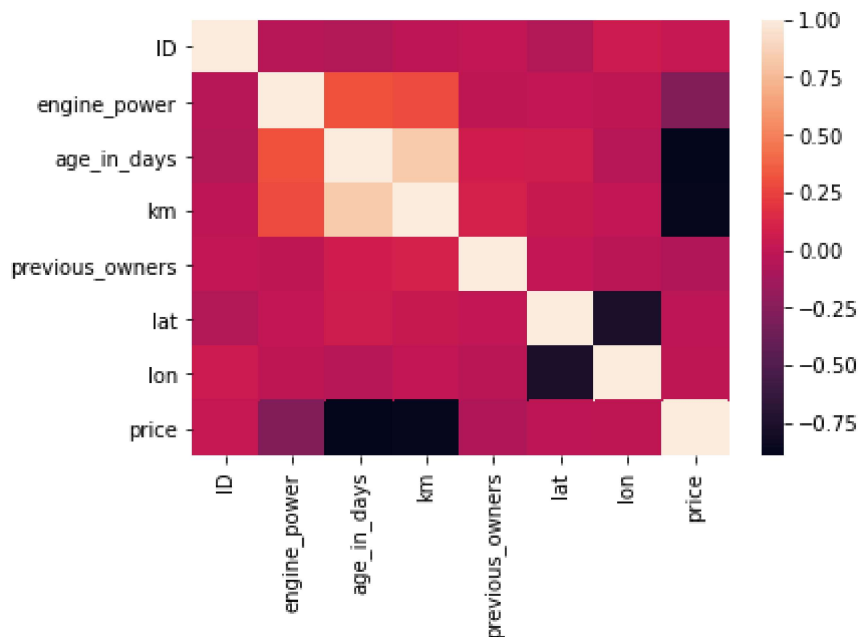
warnings.warn(msg, FutureWarning)

```
Out[9]: <AxesSubplot:xlabel='price', ylabel='Density'>
```



```
In [10]: d1=d[['ID','engine_power', 'age_in_days', 'km', 'previous_owners',
              'lat', 'lon', 'price']]
sns.heatmap(d1.corr())
```

Out[10]: <AxesSubplot:>



```
In [11]: x=d1[['ID','engine_power', 'age_in_days', 'km', 'previous_owners','lat', 'lon']]
y =d1['price']
```

```
In [12]: from sklearn.model_selection import train_test_split
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=0.3)
```

```
In [13]: from sklearn.linear_model import LinearRegression
```

```
In [14]: lr=LinearRegression()
lr.fit(x_train,y_train)
```

Out[14]: LinearRegression()


In [15]: `print(lr.intercept_)`

8817.24210723314

In [ ]:

In [16]: `prediction =lr.predict(x_test)`  
`py.scatter(y_test,prediction)`

Out[16]: <matplotlib.collections.PathCollection at 0x1f3d0ec6250>



In [17]: `print(lr.score(x_test,y_test))`

0.8204307134941244

In [18]: `print(lr.score(x_train,y_train))`

0.8510592128642411

In [19]: `from sklearn.linear_model import Ridge,Lasso`

In [21]: `rr=Ridge(alpha=10)`  
`rr.fit(x_train,y_train)`

Out[21]: Ridge(alpha=10)

In [22]: `rr.score(x_test,y_test)`

Out[22]: 0.8204305148861909

In [23]: `la=Lasso(alpha=10)`  
`la.fit(x_train,y_train)`

Out[23]: **Lasso(alpha=10)**

In [24]: `la.score(x_test,y_test)`

Out[24]: **0.8203904923559537**

In [ ]: